

Kubernetes Orchestration Tools Q&A

In-Class Assignment 2

September 30, 2025

1 Question 1: Key Role of Orchestration Tools

Question: Orchestration tools, such as Kubernetes, play a key role in the server infrastructure for the modern applications.

- (a) Explain how these tools help manage and scale application servers.
- (b) Describe how orchestration tools facilitate automated deployment, scaling, and management of application servers.

Answer:

1.1 (a) How orchestration tools help manage and scale application servers

- **Resource Management:** Automatically allocate and manage computing resources (CPU, memory, storage), ensuring applications get the resources they need
- **Load Balancing:** Distribute traffic across multiple server instances to improve performance and availability
- **Health Monitoring:** Continuously monitor application status and automatically restart failed instances
- **Horizontal Scaling:** Automatically increase or decrease the number of server instances based on load
- **Service Discovery:** Automatically manage network communication and dependencies between services

1.2 (b) How orchestration tools facilitate automated deployment, scaling, and management

- **Declarative Configuration:** Define desired state through YAML files, and the system automatically maintains that state
- **Rolling Updates:** Update application versions without downtime
- **Auto-scaling:** Automatically adjust the number of instances based on CPU usage, memory, or custom metrics
- **Fault Recovery:** Automatically detect and replace failed nodes or containers

- **Configuration Management:** Centrally manage application configurations and secrets

2 Question 2: Difference between Pod, Deployment, and Service

Question: Explain the difference between a Pod, Deployment, and Service.

Answer:

- **Pod:**
 - The smallest deployable unit in Kubernetes
 - Contains one or more tightly coupled containers
 - Shares network and storage resources
 - Has a short lifecycle and can be deleted and recreated at any time
- **Deployment:**
 - A higher-level controller that manages Pod replicas
 - Ensures a specified number of Pod replicas are always running
 - Supports rolling updates and rollbacks
 - Provides declarative application deployment and management
- **Service:**
 - Provides a stable network access point for Pods
 - Routes traffic to backend Pods through label selectors
 - Provides load balancing functionality
 - Maintains consistent IP and DNS names even when Pods are recreated

3 Question 3: Namespace in Kubernetes

Question: What is a Namespace in Kubernetes? Please list one example.

Answer:

A **Namespace** is a virtual cluster concept in Kubernetes used to create multiple logically isolated environments within the same physical cluster. It provides:

- Resource isolation and organization
- Access control and permission management
- Resource quota limitations
- Unique resource names within the namespace

Example:

```
apiVersion: v1
kind: Namespace
metadata:
  name: development
```

Common default namespaces include:

- **default:** Default namespace
- **kube-system:** System components namespace
- **kube-public:** Public resources namespace

4 Question 4: Role of Kubelet and Node Checking Commands

Question: Explain the role of the Kubelet. How do you check the nodes in a Kubernetes cluster? (kubectl command expected)

Answer:

4.1 Role of Kubelet

- Primary agent running on each node
- Manages Pod lifecycle on the node
- Communicates with the API Server to receive Pod specifications
- Monitors container health and reports to the control plane
- Manages container startup, shutdown, and restart
- Performs health checks and resource monitoring

4.2 kubectl commands to check cluster nodes

```
# View all nodes
kubectl get nodes

# View detailed node information
kubectl get nodes -o wide

# View detailed description of a specific node
kubectl describe node <node-name>

# View node status with labels
kubectl get nodes --show-labels
```

5 Question 5: Difference between ClusterIP, NodePort, and LoadBalancer Services

Question: What is the difference between ClusterIP, NodePort, and LoadBalancer services?

Answer:

- **ClusterIP:**
 - Default service type
 - Only accessible within the cluster
 - Assigns an internal cluster IP address
 - Suitable for internal service communication
- **NodePort:**
 - Opens a specific port on each node (30000-32767)
 - Accessible via any node's IP:NodePort
 - Automatically creates ClusterIP
 - Suitable for development and testing environments
- **LoadBalancer:**
 - Creates an external load balancer (requires cloud provider support)
 - Automatically creates NodePort and ClusterIP
 - Provides an externally accessible IP address
 - Suitable for production external access

6 Question 6: Scaling Deployment to 5 Replicas using kubectl

Question: How do you scale a Deployment to 5 replicas using kubectl?

Answer:

```
# Method 1: Using scale command
kubectl scale deployment <deployment-name> --replicas=5

# Method 2: Using patch command
kubectl patch deployment <deployment-name> -p '{"spec":{"replicas":5}}'

# Verify scaling result
kubectl get deployment <deployment-name>
```

7 Question 7: Updating Deployment Image without Downtime

Question: How would you update the image of a Deployment without downtime?

Answer:

```
# Method 1: Using set image command
kubectl set image deployment/<deployment-name> <container-name>=<new-image>

# Method 2: Using patch command
kubectl patch deployment <deployment-name> -p '{"spec":{"template":{"spec":{"containers":[{"name":"<container-name>","image":"<new-image>"}]}}}}'

# Check rolling update status
kubectl rollout status deployment/<deployment-name>

# Rollback if needed
kubectl rollout undo deployment/<deployment-name>
```

8 Question 8: Exposing Deployment to External Traffic

Question: How do you expose a Deployment to external traffic?

Answer:

```
# Method 1: Create NodePort service
kubectl expose deployment <deployment-name> --type=NodePort --port=80

# Method 2: Create LoadBalancer service
kubectl expose deployment <deployment-name> --type=LoadBalancer --port=80

# Method 3: Using kubectl create service
kubectl create service nodeport <service-name> --tcp=80:8080

# Method 4: Through Ingress (need to create Service first)
kubectl expose deployment <deployment-name> --port=80
# Then create Ingress resource
```

9 Question 9: Kubernetes Scheduling Decision Process

Question: How does Kubernetes scheduling decide which node a Pod runs on?

Answer:

The Kubernetes scheduler decides which node a Pod runs on through the following steps:

9.1 Scheduling Process

1. **Filtering Phase:** Exclude nodes that don't meet requirements
 - Resource requirements (CPU, memory)
 - Node selectors and affinity rules
 - Taints and tolerations
 - Port conflict checks
2. **Scoring Phase:** Score eligible nodes
 - Resource utilization balance
 - Affinity preferences
 - Image locality
 - Node load balancing
3. **Selection Phase:** Choose the highest-scoring node

9.2 Influencing Factors

- Node resource availability
- Pod resource requests and limits
- Node selectors (nodeSelector)
- Affinity and anti-affinity rules
- Taints and Tolerations

10 Question 10: Role of Ingress and Difference from Service

Question: What is the role of Ingress and how does it differ from a Service?

Answer:

10.1 Role of Ingress

- Manages external HTTP/HTTPS access to services in the cluster
- Provides load balancing, SSL termination, and name-based virtual hosting
- Supports path-based and domain-based routing
- Centrally manages external access rules

10.2 Difference between Ingress and Service

Feature	Service	Ingress
Layer	L4 (Transport Layer)	L7 (Application Layer)
Protocol	TCP/UDP	HTTP/HTTPS
Routing	Port-based	Path/domain-based
SSL	Not supported	Supports SSL termination
Load Balancing	Simple round-robin	Advanced load balancing
Cost	Each service needs Load-Balancer	Single entry point

10.3 Example Configuration

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
spec:
  rules:
  - host: example.com
    http:
      paths:
      - path: /app1
        pathType: Prefix
        backend:
          service:
            name: app1-service
            port:
              number: 80

```